

CSCE 435 Group project

0. Group number: 8

1. Group members:

1. Alex Do
2. Alex Byrd
3. Jose Rojo
4. Matthew Livesay

1.1 Communication:

We will be communicating using an iMessage group chat. This has been created already and all members have responded.

2. Project topic (e.g., parallel sorting algorithms)

The project includes parallelizing sequential sorting algorithms that include bitonic sort, sample sort, merge sort, and radix sort. After parallelizing the algorithms, we will examine their performance by varying the number of input sizes, the number of processors involved in the operation, and how the initial input array is generated.

2a. Brief project description (what algorithms will you be comparing and on what architectures)

- Bitonic Sort (Alex Do): Bitonic Sort is a divide-and-conquer algorithm that operates by constructing a sequence of elements that forms a bitonic sequence, which is basically a sequence that first increases and then decreases. The algorithm then recursively sorts this bitonic sequence by performing compare-exchange operations to produce a sorted sequence. When bitonic sort is parallelized, the core operations of comparing and exchanging elements are distributed across multiple processors.
- Sample Sort (Alex Byrd): The Parallel Sample Sort algorithm distributes data across multiple processes, where each process (including the master) sorts its chunk concurrently using some sequential sorting method (in this case, quicksort). After sorting, processes send samples to the master, which selects splitters and broadcasts them. Each process splits its sorted chunk into buckets based on the splitters, exchanges buckets with other processes, and sorts them locally. The final sorted data is then gathered and merged, efficiently balancing computation and communication across all processes.
- Merge Sort (Jose Rojo): Merge sort is a divide-and-conquer sorting algorithm that recursively splits an array into two halves, sorts each half, and then merges the sorted halves back together. The process continues until the array is split into individual elements, which are inherently sorted. Then, during the merging phase, the sorted subarrays are combined to produce a fully sorted array.

- Radix Sort (Matthew Livesay): Radix sort works by sorting an array from LSB to MSB. A group of bits is taken into account and then the entire array is sorted to make the considered bits ordered from smallest to largest. By the time the algorithm has completed sorting the MSBs, the entire array will be sorted.

2b. Pseudocode for each parallel algorithm

- For MPI programs, include MPI calls you will use to coordinate between processes

Bitonic Sort

```
function ParallelBitonicSort()
    Initialize MPI w/ MPI_Init()
    Get rank and size w/ MPI_Comm_rank() and MPI_Comm_size()

    total_elements = Get from user input and verify it's 2^n
    elements_per_process = total_elements / size

    // Scattering input array
    if rank == 0
        global_array = initialize_array(total_elements)
    else
        global_array = None

    local_array = Allocate array of size elements_per_process
    Scatter portions of the global array from the root process w/
    MPI_Scatter()

    // Main bitonic sort loop
    for k = 2 to total_elements by multiplying by 2
        for j = k // 2 down to 1 by dividing by 2

            // Determine sorting direction
            if (rank & (k // 2)) == 0
                ascending = true
            else
                ascending = false

            // Calculate partner process
            partner = rank XOR j

            if partner < size
                // Perform exchange and merge
                CompareExchange(local_array, partner, ascending)

    Gather data from all processes and assembles it into a single array on
    the root process w/ MPI.Gather()

    Finalize MPI w/ MPI.Finalize()

    // Helper functions during bitonic sort
    function compareExchange(int local_array[], int partner, bool ascending)
```

```

    if rank < partner
        Send local_array to partner
        Receive partner_array from partner
    else
        Receive partner_array from partner
        Send local_array to partner

    combined_array = Merge(local_array, partner_array, ascending)

    // Determine which half to keep
    if ( (rank < partner and ascending) or (rank > partner and not
ascending) )
        local_array = first half of combined_array
    else
        local_array = second half of combined_array

function merge(int array1[], int array2[], bool ascending)
    merged_array = array1 + array2
    Sort merged_array in ascending or descending order based on
'ascending' flag
    return merged_array

```

Radix Sort

```

each processor starts with a portion of the array to be sorted (will
likely generate it)

for each chunk of bits
    iterate over entire portion of array to generate a histogram

    send histogram to all other processes
    receive histograms from all other processes using mpi_reduce or
similar

    combine all histograms (this might be accomplished by the mpi call that
collects all histograms)
    calculate prefix sum array using histogram array

    calculate processor offset using mpi_rank
    combine offset and prefix sum array to find final offset for each
value in array portion

    use MPI call to place each value in a global data structure

```

Merge Sort

```

// Master process
    for i = 1 to array_size - 1
        Use MPI Send to even amount of data to all worker processes

    for i = 1 to size - 1
        Use MPI Recieve to get sorted arrays from worker processes
        call "merge" function to merge sorted arrays

// Worker processes
    Use MPI_recv to recieve arrays from master process
    call "sequentialMergeSort" function to sort recieved array
    Use MPI_send sorted arrays back to the master process

//Sequential Merge Sort
function mergeSort(array A, int left, int right)
    if (left < right) // Check if the array has more than one element
        int mid = (left + right) / 2

        // Sort the left half
        mergeSort(A, left, mid)

        // Sort the right half
        mergeSort(A, mid + 1, right)

        // Merge the sorted halves
        merge(A, left, mid, right)
end function

function merge(array A, int left, int mid, int right)
    // Create temporary arrays to hold the left and right halves
    int leftArray[mid - left + 1]
    int rightArray[right - mid]

    // Copy data to temporary arrays
    for i = 0 to mid - left
        leftArray[i] = A[left + i]
    for j = 0 to right - mid - 1
        rightArray[j] = A[mid + 1 + j]

    // Merge the temporary arrays back into A
    int i = 0, j = 0, k = left
    while (i < size of leftArray and j < size of rightArray)
        if (leftArray[i] <= rightArray[j])
            A[k] = leftArray[i]
            i++
        else
            A[k] = rightArray[j]
            j++
        k++

    // Copy remaining elements of leftArray, if any
    while (i < size of leftArray)
        A[k] = leftArray[i]

```

```

        i++
        k++

    // Copy remaining elements of rightArray, if any
    while (j < size of rightArray)
        A[k] = rightArray[j]
        j++
        k++
end function

```

Sample Sort

Initialize MPI (MPI_init, MPI_Comm size & rank)

Note: The master process basically has the worker process tasks weaved into it, as it only does additional computations when the workers are done. If we didn't treat the master process as an additional worker process, we would lose a good amount of performance.

```

// Master process
    for (worker processes)
        Use MPI_Send to even amount of data to all worker processes

    Sort the master's chunk

    MPI Gather the samples from each processor
    Sort samples (use quicksort here, only master process)
    Select (m-1) splitters
    MPI Broadcast splitters to all processors

    Split the chunk into buckets based on splitters

    Gather sorted buckets from all processes
    for (worker processes)
        recv data for each process

    Merge the sorted buckets

// Worker processes
    Use MPI_recv to recieve arrays from master process
    Sort the chunks (use quicksort here, per process)
    Gather sampled elements back to the master
    Recv splitters from master
    Split the chunk into buckets based on splitters
    Send the corresponding buckets back

```

The Parallel Sample Sort algorithm distributes data across multiple processes, where each process (including the master) sorts its chunk concurrently using some sequential sorting method (in this case,

quicksort). After sorting, processes send samples to the master, which selects splitters and broadcasts them. Each process splits its sorted chunk into buckets based on the splitters, exchanges buckets with other processes, and sorts them locally. The final sorted data is then gathered and merged, efficiently balancing computation and communication across all processes.

2c. Evaluation plan - what and how will you measure and compare

- Input sizes, Input types
 - The input array sizes will always be 2^N , and therefore of length 2^{16} , 2^{18} , 2^{20} , 2^{22} , 2^{24} , 2^{26} , 2^{28}
 - These input array types will be either sorted, random, reverse sorted, or 1% perturbed
 - All values in the input arrays will be integers
- Strong scaling (same problem size, increase number of processors/nodes)
 - We will analyze how the sorting algorithms scale when increasing the number of processors while keeping the problem size constant, allowing us to determine how well an algorithm can take advantage of additional computational resources for the same problem. That is, the execution time should decrease as more processors are added.
 - Therefore, we will compare the execution time of the parallel sorting algorithms with varying processor counts (2, 4, 8, 16, 32, 64, 128, 256, 512, 1024) for each input array size
- Weak scaling (increase problem size, increase number of processors)
 - We will evaluate the performance when both the problem size and the number of processors increase proportionally, allowing us to determine if the algorithm can handle larger problems as more resources (processors) are added. That is, the algorithm should be maintaining a constant execution time as the processors and problem size scale together
 - Therefore, we will compare the execution time of the parallel sorting algorithms with increasing array sizes (2^{16} , 2^{18} , 2^{20} , 2^{22} , 2^{24} , 2^{26} , 2^{28}) with corresponding increasing processor counts (2, 4, 8, 16, 32, 64, 128, 256, 512, 1024)